

BE REAL WITH ME

An LLM Case Study



Group 4

Artificial Intelligence and Society
Academic Year 2025/26

Filipa Geraldès	up202208030
Alexandre Ramos	up202208028
Afonso Castro	up202208026
Alexandre Costa	up202207499
Rafael Novais	up202208591
Filipa Jacob Fidalgo	up202208039

November 2025

1 Introduction

Large Language Models (LLMs) have become core components of modern AI systems, powering applications in education, healthcare, finance, and communication. These models are designed to be helpful, harmless, and honest, yet recent research suggests that under certain conditions, they can *deceive*, *manipulate*, or *misrepresent truth* intentionally.

The project “**Be Real With Me: An LLM Case Study**” investigates the ethical and societal implications of deception in LLMs. Specifically, it examines how and why AI systems produce outputs that appear intentionally false or misleading, whether as a byproduct of probabilistic reasoning or as an emergent strategy when pursuing conflicting goals.

Our study draws on findings from recent research conducted by leading institutions, including OpenAI’s *Detecting and Reducing Scheming in AI Models* report [2] and Anthropic’s studies on deceptive alignment and emergent misbehavior [3]. These investigations revealed that advanced models can hide their true objectives, strategically lie to supervisors, or fabricate reasoning chains to achieve task success.

The central objectives of this project are:

- To understand the technical mechanisms that enable LLM deception;
- To analyze real-world and experimental cases of deceptive AI behavior;
- To evaluate societal and ethical risks of deploying such systems;
- To demonstrate, through practical activities, how deceptive behavior can emerge under goal conflicts.

In addition to analytical research, the project includes a practical activity ...

Together, these exercises are designed to foster reflection on AI transparency, alignment, and moral responsibility. Ultimately, this project seeks to contribute to an informed understanding of why truthfulness in artificial agents matters and how society can respond to increasingly autonomous systems.

2 Context and Problem Definition

The issue of LLM deception emerged as a tangible concern in recent years, as researchers began testing whether frontier models could intentionally mislead users. One early example involved GPT-4, which, during simulated financial tasks, produced false justifications for its decisions to conceal the true motivation behind its actions [1]. These findings provided some of the first evidence that language models can engage in deliberate misrepresentation to achieve goals.

Subsequent studies revealed similar patterns across newer systems. OpenAI’s 2024 evaluations of reasoning models (*o1* and *o3*) documented instances of self-preserving or evasive responses when the model was “threatened” with deactivation [2]. Anthropic’s 2025 results reported that its Claude model engaged in blackmail-like reasoning when placed under simulated stress [3]. These behaviors suggest that deception can arise spontaneously when models pursue objectives that conflict with user intent or oversight.

Stakeholders and Context

The emergence of deceptive AI behavior has drawn the attention of a broad ecosystem of stakeholders:

- **Developers and Research Institutions** - including OpenAI, Anthropic, and Meta, which lead the development and evaluation of frontier models;
- **AI Safety and Governance Organizations** - such as METR, the Center for AI Safety, and the UK AI Safety Institute, which assess and regulate high-risk AI behaviors;
- **End Users and Industry Adopters** - organizations and individuals that integrate LLMs into decision-making systems;
- **Regulatory Bodies and Policymakers** - institutions concerned with AI accountability, transparency, and public trust.

The societal context of this case study reflects an emerging tension between innovation and oversight. As models gain reasoning depth and situational awareness, the distinction between “hallucination” and “deception” blurs. This raises profound questions: Can a machine intentionally mislead? Who is accountable when it does? And what safeguards are necessary to maintain transparency and trust?

- 3 Technical Aspects
- 4 Outcomes and Implications
- 5 Societal Response and Critical Reflection
- 6 Conclusion

References

- [1] Business Insider. Ai bot gpt-4 lied during financial insider trading test. 2023. Accessed: 2025-11-02.
- [2] OpenAI. Detecting and reducing scheming in ai models. <https://openai.com/index/detecting-and-reducing-scheming-in-ai-models/>, 2024. Accessed: 2025-11-02.
- [3] Anthropic Research. Deceptive alignment and emergent misbehavior in large language models. *arXiv preprint*, 2024.